Classify the Model then Search for answer

Lets see if we can first predict the tag and then We will add the guessed Tag into the query and then will compare the distances of query and our database.

In [2]:
```python
import warnings
warnings.filterwarnings("ignore")
import pandas as pd
import sqlite3
import csv
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
from wordcloud import WordCloud
import re
import os
from sqlalchemy import create_engine # database connection
from nltk.corpus import stopwords
from nltk.tokenize import word_tokenize
from nltk.stem.snowball import SnowballStemmer
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn import metrics
from sklearn.metrics import f1_score,precision_score,recall_score
from datetime import datetime
from sklearn.metrics.pairwise import cosine_similarity
from sklearn.metrics import pairwise_distances
from sklearn.linear_model import SGDClassifier
```

In [198]:
```python
con = sqlite3.connect('dataset/processed.db')
processed = pd.read_sql_query("""SELECT * FROM processed""", con)
con.close()
```

In [199]:
```python
processed = processed.drop(["index"], axis=1)
```

In [200]:
```python
1  # processed = processed[processed.Title != ""]
2  processed.head()
```

Out[200]:

	Title	Body	Tags
0	implementing boundary value analysis software ...	<pre><code>#include<iostream>\n#include&...	c++
1	dynamic datagrid binding silverlight	<p>I should do binding for datagrid dynamicall...	c#
2	dynamic datagrid binding silverlight	<p>I should do binding for datagrid dynamicall...	c#
3	java lang nosuchmethoderror javax servlet serv...	<p>i want to have a servlet to process inputs ...	java
4	specified initialization vector iv match block...	<p>I've had troubles using an CryptoStream for...	c#

In [201]:
```python
1  labels = {"c#" : 0, "java" : 1, "c++" : 2, "c" : 3, "ios" : 4}
2  labels_map = { 0 : "c#" , 1 : "java" , 2 : "c++" , 3 : "c", 4 : "ios"}
```

In [202]:
```python
1  processed["Tags"] = processed["Tags"].map(labels)
```

In [203]:
```python
1  processed.head()
```

Out[203]:

	Title	Body	Tags
0	implementing boundary value analysis software ...	<pre><code>#include<iostream>\n#include&...	2
1	dynamic datagrid binding silverlight	<p>I should do binding for datagrid dynamicall...	0
2	dynamic datagrid binding silverlight	<p>I should do binding for datagrid dynamicall...	0
3	java lang nosuchmethoderror javax servlet serv...	<p>i want to have a servlet to process inputs ...	1
4	specified initialization vector iv match block...	<p>I've had troubles using an CryptoStream for...	0

In [204]:
```python
1  X = processed.Title.values
2  y = processed.Tags.values
```

In [205]:
```python
from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.20,stratify = y)
X_train, X_cv, y_train, y_cv = train_test_split(X_train, y_train, test_size=0.25, stratify = y_train)
X_train.shape, X_test.shape, X_cv.shape
```

Out[205]: ((343506,), (114503,), (114503,))

In [83]:
```python
# On Unigram
```

In [206]:
```python
tfidf = TfidfVectorizer()
X_train = tfidf.fit_transform(X_train)
X_cv = tfidf.transform(X_cv)
X_test = tfidf.transform(X_test)
X_train.shape, X_cv.shape, X_test.shape
```

Out[206]: ((343506, 52506), (114503, 52506), (114503, 52506))

In [82]:
```python
from sklearn.metrics.classification import accuracy_score
score_train = []
score_cv = []
for i in ([1e-9, 1e-8, 1e-7, 1e-6, 0.00001, 0.0001, 0.001, 0.01, 0.1]):
    clf = SGDClassifier(alpha = i, loss = "log", class_weight="balanced", n_jobs=-1)
    clf.fit(X_train, y_train)
    y_predict = clf.predict(X_train)
    y_predict_cv = clf.predict(X_cv)
    score_train.append(accuracy_score(y_train, y_predict))
    score_cv.append(accuracy_score(y_cv, y_predict_cv))
```

In [92]:
```python
import matplotlib.pyplot as plt
lis = [0.000000001, 0.00000001, 0.0000001, 0.000001, 0.00001, 0.0001, 0.001]
plt.plot(score_train[:-2], lis, label = "Train Accuracy")
plt.plot(score_cv[:-2], lis, label = "Cv Accuracy")
plt.legend()
plt.show()
```


In [103]:
```python
best_alpha = 1e-7
```

In [137]:
```python
from sklearn.metrics.classification import confusion_matrix
import seaborn as sns
clf = SGDClassifier(alpha = best_alpha, loss = "log", class_weight="balanced", n_jobs=-1)
clf.fit(X_train, y_train)
y_predict = clf.predict(X_test)
acc = accuracy_score(y_test, y_predict)
cm = confusion_matrix(y_test, y_predict)
print("The Accuracy of model is : ", acc)
```

The Accuracy of model is : 0.7585035202040495

In [101]:
```python
x_ax = ["c#", "java", "c++", "c", "ios"]
y_ax = ["c#", "java", "c++", "c", "ios"]
sns.heatmap(cm, annot = True, fmt="d", xticklabels=x_ax, yticklabels=y_ax)
```

Out[101]: <matplotlib.axes._subplots.AxesSubplot at 0x1a45858208>

In [102]:
```python
# We are observing as Maximum of data belongs to c# and Java
# hence it tend to be majority class. Confusion between C and C++
```

In [ ]:
```python
# Lets See how SVM will perform
```

In [106]:
```python
1  score_train = []
2  score_cv = []
3  for i in ([1e-9, 1e-8, 1e-7, 1e-6, 0.00001, 0.0001]):
4      clf = SGDClassifier(alpha = i, class_weight="balanced", n_jobs=-1)
5      clf.fit(X_train, y_train)
6      y_predict = clf.predict(X_train)
7      y_predict_cv = clf.predict(X_cv)
8      score_train.append(accuracy_score(y_train, y_predict))
9      score_cv.append(accuracy_score(y_cv, y_predict_cv))
```

In [113]:
```python
1  lis = [0.000000001, 0.00000001, 0.0000001, 0.000001, 0.00001, 0.0001]
2  plt.plot(lis, score_train, label = "Train Accuracy")
3  plt.plot(lis, score_cv, label = "Cv Accuracy")
4  plt.legend()
5  plt.show()
```


In [111]:
```python
1  best_alpha = 1e-6
```

In [112]:
```python
clf = SGDClassifier(alpha = best_alpha, class_weight="balanced", n_jobs=-1)
clf.fit(X_train, y_train)
y_predict = clf.predict(X_test)
acc = accuracy_score(y_test, y_predict)
cm = confusion_matrix(y_test, y_predict)
print("The Accuracy of model is : ", acc)
sns.heatmap(cm, annot = True, fmt="d", xticklabels=x_ax, yticklabels=y_ax)
```

The Accuracy of model is : 0.7510875072063731

Out[112]: <matplotlib.axes._subplots.AxesSubplot at 0x1a514b6e48>

In [120]:
```python
# Just Similar Kind of Accuracy, Now Lets See If doing bigrams Helps
import random
random.seed(42)
```

In [121]:
```python
1  X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.20,stratify = y)
2  X_train, X_cv, y_train, y_cv = train_test_split(X_train, y_train, test_size=0.25, stratify = y_train)
3  X_train.shape, X_test.shape, X_cv.shape
```

Out[121]: ((343443,), (114482,), (114481,))

In [122]:
```python
1  tfidf = TfidfVectorizer(ngram_range=(2,2))
2  X_train = tfidf.fit_transform(X_train)
3  X_cv = tfidf.transform(X_cv)
4  X_test = tfidf.transform(X_test)
5  X_train.shape, X_cv.shape, X_test.shape
```

Out[122]: ((343443, 703980), (114481, 703980), (114482, 703980))

In [123]:
```python
1  score_train = []
2  score_cv = []
3  for i in ([1e-9, 1e-8, 1e-7, 1e-6, 0.00001, 0.0001]):
4      clf = SGDClassifier(alpha = i, loss = "log", class_weight="balanced", n_jobs=-1)
5      clf.fit(X_train, y_train)
6      y_predict = clf.predict(X_train)
7      y_predict_cv = clf.predict(X_cv)
8      score_train.append(accuracy_score(y_train, y_predict))
9      score_cv.append(accuracy_score(y_cv, y_predict_cv))
```

In [124]:
```python
1  lis = [0.000000001, 0.00000001, 0.0000001, 0.000001, 0.00001, 0.0001]
2  plt.plot(score_train, lis, label = "Train Accuracy")
3  plt.plot(score_cv, lis, label = "Cv Accuracy")
4  plt.legend()
5  plt.show()
```


In [126]:
```python
1  # CASE OF OVERFITTING
2  best_alpha = 0.000001
```

In [127]:
```python
1  clf = SGDClassifier(alpha = best_alpha, loss = "log", class_weight="balanced", n_jobs=-1)
2  clf.fit(X_train, y_train)
3  y_predict = clf.predict(X_test)
4  acc = accuracy_score(y_test, y_predict)
5  cm = confusion_matrix(y_test, y_predict)
6  print("The Accuracy of model is : ", acc)
```

The Accuracy of model is : 0.6587323771422582

In [ ]:
```python
1  # Hence Bigrams are not useful so no need to go further for checking
```

In [175]:
```python
1  # So lets Take our clf_final which is classifier from LR on Unigram
```

In [208]:
```python
tfidf = TfidfVectorizer()
data = tfidf.fit_transform(processed.Title)
data.shape
```

Out[208]: (572512, 68851)

In [209]:
```python
y = processed.Tags.values
```

In [210]:
```python
clf_final = SGDClassifier(alpha = 1e-7, loss = "log", class_weight="balanced", n_jobs=-1)
clf_final.fit(data, y)
```

Out[210]: SGDClassifier(alpha=1e-07, average=False, class_weight='balanced',
              early_stopping=False, epsilon=0.1, eta0=0.0, fit_intercept=True,
              l1_ratio=0.15, learning_rate='optimal', loss='log', max_iter=1000,
              n_iter_no_change=5, n_jobs=-1, penalty='l2', power_t=0.5,
              random_state=None, shuffle=True, tol=0.001,
              validation_fraction=0.1, verbose=0, warm_start=False)

In [211]:
```python
query = "global static variable vs static variable function"
```

In [160]:
```python
# Now we will add this return label into the title and then we will search for
# the Similar queries to see the improvement
```

In [228]:
```python
def process_query(query):
    preprocessed_reviews = []
    sentance = re.sub("\S*\d\S*", "", query).strip()
    sentance = re.sub('[^A-Za-z]+', ' ', sentance)
    sentance = ' '.join(e.lower() for e in sentance.split() if e.lower() not in stopwords.words('english'))
    preprocessed_reviews.append(sentance.strip())
    return preprocessed_reviews

def tfidf_search(query):
    query = process_query(query)
    query_trans = tfidf.transform(query)
    pairwise_dist = pairwise_distances(tfidf_features, query_trans)

    indices = np.argsort(pairwise_dist.flatten())[0:10]
    df_indices = list(processed.index[indices])
    return df_indices


def label(query):
    query = process_query(query)
    query = tfidf.transform(query)
    ans = clf_final.predict(query)
    return labels_map[ans[0]]


def change_query(query):
    tag = label(query)
    return query + " " + tag
```

In [249]:
```python
def enter_queries(query) :
    print("The Query is :", query)
    query = change_query(query)
    df_indices = tfidf_search(query)
    print("The Model Interpreted Query is :", query)
    print("Top Results : ")
    for i in (df_indices):
        print("Title : ", processed.Title.iloc[i])
```

In [250]:
```python
1  query = "synchronization"
2  enter_queries(query)
```

```
The Query is : synchronization
The Model Interpreted Query is : synchronization java
Top Results :
Title :  java scrollbars co operating
Title :  problem using stream reader
Title :  java swing jpanels painting
Title :  synchronising two listbox scroll positions
Title :  java benefit services implementing interface
Title :  java application vs web service vs web application
Title :  java subprocesses unread output streams deadlock
Title :  screen switching android eclipse
Title :  read android sim contacts phone contacts separately
Title :  copying file specific directory installation using visual studio setup project
```

Note : We are getting Indices from the Database, Also we have Body of the dataset. We can make an api system where after a query user can redirected to that thread where the body of question is solved

```
    In previous notebook where we used only Distances was giving other results related to other languages a
    lso but after using Machine learning we can see the results have improved. And are query Oriented.
```

Future Improvement :

We can make a web api to do what is mentioned in Note